JOB SHOP SCHEDULING BY SIMULATED ANNEALING

PETER J. M. VAN LAARHOVEN
Nederlandse Philips Bedrijven, Eindhoven, The Netherlands

EMILE H. L. AARTS
Philips Research Laboratories, Eindhoven, The Netherlands
and
Eindhoven University of Technology, Eindhoven, The Netherlands

JAN KAREL LENSTRA
Eindhoven University of Technology, Eindhoven, The Netherlands and CWI, Amsterdam, The Netherlands

We describe an approximation algorithm for the problem of finding the minimum makespan in a job shop. The algorithm is based on simulated annealing, a generalization of the well known iterative improvement approach to combinatorial optimization problems. The generalization involves the acceptance of cost-increasing transitions with a nonzero probability to avoid getting stuck in local minima. We prove that our algorithm asymptotically converges in probability to a globally minimal solution, despite the fact that the Markov chains generated by the algorithm are generally not irreducible. Computational experiments show that our algorithm can find shorter makespans than two recent approximation approaches that are more tailored to the job shop scheduling problem. This is, however, at the cost of large running times.

W e are concerned with a problem in machine scheduling known as the *job shop scheduling problem* (Coffman 1976, French 1982). Informally, the problem can be described as follows. We are given a set of jobs and a set of machines. Each job consists of a chain of operations, each of which needs to be processed during an uninterrupted time period of a given length on a given machine. Each machine can process at most one operation at a time. A *schedule* is an allocation of the operations to time intervals on the machines. The problem is to find a schedule of minimum length.

The job shop scheduling problem is among the hardest combinatorial optimization problems. Not only is it NP-hard, but even among the members of the latter class it appears to belong to the more difficult ones (Lawler, Lenstra and Rinnooy Kan 1982). *Optimization algorithms* for job shop scheduling proceed by branch and bound, see, for instance, Lageweg, Lenstra and Rinnooy Kan (1977), and Carlier and Pinson (1989). Most *approximation algorithms* use a *priority rule*, i.e., a rule for choosing an operation from a specified subset of as yet unscheduled opera-

tions. Adams, Balas and Zawack (1988) developed a *shifting bottleneck procedure*, which employs an ingenious combination of schedule construction and iterative improvement, guided by solutions to single-machine problems. The approach pursued by Adams, Balas and Zawack is strongly tailored to the problem at hand. It is based on a fair amount of combinatorial insight into the job shop scheduling problem, and its implementation requires a certain level of programmer sophistication.

In this paper, we investigate the potential of a more general approach based on the easily implementable *simulated annealing algorithm* (Kirkpatrick, Gelatt and Vecchi 1983, and Černy 1985). A similar approach was independently investigated by Matsuo, Suh and Sullivan (1988). Their *controlled search simulated annealing* method is less general than ours in the sense that it uses more problem specific neighborhoods.

The organization of this paper is as follows. In Section 1 we give a formal problem definition. In Section 2 the basic elements of the simulated annealing algorithm are reviewed. In Section 3 we describe

the application of simulated annealing to job shop scheduling. We prove asymptotic convergence of the algorithm to a globally minimal solution by showing that the neighborhood structure is such that each ergodic set contains at least one global minimum. Section 4 contains the results of a computational study in which simulated annealing is used to find approximate solutions to a large set of instances of the job shop scheduling problem. We compare our simulated annealing method with three other approaches, i.e., time-equivalent iterative improvement, the shifting bottleneck procedure of Adams, Balas and Zawack, and the controlled search simulated annealing method of Matsuo, Suh and Sullivan. In Section 5 we end with some concluding remarks.

1. THE PROBLEM

We are given a set $\mathcal{J}$ of n jobs, a set $\mathcal{M}$ of m machines, and a set $\mathcal{O}$ of N operations. For each operation $v \in \mathcal{O}$ there is a job $J_v \in \mathcal{J}$ to which it belongs, a machine $M_v \in \mathcal{M}$ on which it requires processing, and a processing time $t_v \in \mathbb{N}$. There is a binary relation $\rightarrow$ on $\mathcal{O}$ that decomposes $\mathcal{O}$ into chains corresponding to the jobs; more specifically, if $v \rightarrow w$, then $J_v = J_w$, and there is no $x \notin \{v, w\}$ such that $v \rightarrow x$ or $x \rightarrow w$. The problem is to find a start time s_v for each operation $v \in \mathcal{O}$ such that

$$\max_{v \in \mathcal{O}} s_v + t_v, \tag{1}$$

is minimized subject to

$$s_v \geq 0 \quad \text{for all } v \in \mathcal{O} \tag{2}$$

$$s_w - s_v \geq t_v \quad \text{if } v \rightarrow w, \quad v, w \in \mathcal{O} \tag{3}$$

$$s_w - s_v \geq t_v \lor s_v - s_w \geq t_w$$

$$\text{if } M_v = M_w, \quad v, w \in \mathcal{O}. \tag{4}$$

It is useful to represent the problem by the *disjunctive graph* model of Roy and Sussmann (1964). The disjunctive graph $G = (V, A, E)$ is defined as follows:

- $V = \mathcal{O} \cup \{0, N + 1\}$, where 0 and $N + 1$ are two fictitious operations; the weight of a vertex v is given by the processing time t_v ($t_0 = t_{N+1} = 0$).
- $A = \{(v, w) \mid v, w \in \mathcal{O}, v \rightarrow w\} \cup \{(0, w) \mid w \in \mathcal{O}, \nexists v \in \mathcal{O}: v \rightarrow w\} \cup \{(v, N + 1) \mid v \in \mathcal{O}, \nexists w \in \mathcal{O} \ v \rightarrow w\}$. Thus, A contains arcs connecting consecutive operations of the same job, as well as arcs from 0 to the first operation of each job and from the last operation of each job to $N + 1$.
- $E = \{\{v, w\} \mid M_v = M_w\}$. Thus, edges in E connect operations to be processed by the same machine.


Figure 1. The disjunctive graph G of a 3-job 3-machine instance. Operations 1, 5 and 9 are processed by machine 1, operations 2, 4 and 8 by machine 2, and operations 3, 6 and 7 by machine 3; 0 and 10 are the fictitious initial and final operations, respectively. Thick arrows denote arcs in A, dotted lines edges in E.

Figure 1 illustrates the disjunctive graph for a 3-job, 3-machine instance, where each job consists of three operations.

For each pair of operations $v, w \in \mathcal{O}$ with $v \rightarrow w$, condition (3) is represented by an arc (v, w) in A. Similarly, for each pair of operations $v, w \in \mathcal{O}$ with $M_v = M_w$, the disjunctive constraint (4) is represented by an edge $\{v, w\}$ in E, and the two ways to settle the disjunction correspond to the two possible orientations of the edge. There is an obvious one-to-one correspondence between a set of choices in (4) that is overall feasible and an orientation of all the edges in E for which the resulting digraph is acyclic. The objective value (the makespan) of the corresponding solution is given by the length of a longest path in this digraph. Such a set of orientations decomposes $\mathcal{O}$ into chains corresponding to the machines, i.e., it defines for each machine an ordering or permuation of the operations to be processed by that machine. Conversely, a set of machine permutations defines a set of orientations of the edges in E, though not necessarily one which results in an acyclic digraph. Since the longest path in a cyclic digraph has infinite length, we can now rephrase the problem as that of finding a set of machine permutations that minimizes the longest path in the resulting digraph. In Section 3 we use this formulation of the problem to find approximate solutions by simulated annealing.

2. SIMULATED ANNEALING

Ever since its introduction, independently by Kirkpatrick, Gelatt and Vecchi (1983) and Černy

(1985), simulated annealing has been applied to many combinatorial optimization problems in such diverse areas as computer-aided design of integrated circuits, image processing, code design and neural network theory; for a review the reader is referred to Van Laarhoven and Aarts (1987). The algorithm is based on an intriguing combination of ideas from at first sight completely unrelated fields of science, viz. combinatorial optimization and statistical physics. On the one hand, the algorithm can be considered as a generalization of the well known iterative improvement approach to combinatorial optimization problems, on the other hand, it can be viewed as an analogue of an algorithm used in statistical physics for computer simulation of the annealing of a solid to its *ground state*, i.e., the state with minimum energy. In this paper, we mainly restrict ourselves to the first point of view; thus, we first briefly review iterative improvement.

Generally, a combinatorial optimization problem is a tuple $(\mathscr{R}, \mathscr{C})$, where $\mathscr{R}$ is the *set of configurations* or *solutions* of the problem, and $C: \mathscr{R} \to \mathbb{R}$ the *cost function* (Papadimitriou and Steiglitz 1982). To be able to use iterative improvement we need a *neighborhood structure* $\mathscr{N}: \mathscr{R} \to 2^{\mathscr{R}}$; thus, for each configuration i, $\mathscr{N}(i)$ is a subset of configurations, called the *neighborhood* of i. Neighborhoods are usually defined by first choosing a simple type of *transition* to obtain a new configuration from a given one and then defining the neighborhood as the set of configurations that can be obtained from a given configuration in one transition.

Given the set of configurations, a cost function and a neighborhood structure, we can define the iterative improvement algorithm as follows. The algorithm consists of a number of iterations. At the start of each iteration, a configuration i is given and a transition to a configuration $j \in \mathscr{N}(i)$ is generated. If $C(j) < C(i)$, the start configuration in the next iteration is j, otherwise it is i. If $\mathscr{R}$ is finite and if the transitions are generated in some exhaustive enumerative way, then the algorithm terminates by definition in a local minimum. Unfortunately, a local minimum may differ considerably in cost from a global minimum. Simulated annealing can be viewed as an attempt to find *near-optimal* local minima by allowing the acceptance of cost-increasing transitions. More precisely, if i and $j \in \mathscr{N}(i)$ are the two configurations to choose from, then the algorithm continues with configuration j with a probability given by $\min\{1, \exp(-(C(j) - C(i))/c)\}$, where c is a positive *control parameter*, which is gradually decreased during the execution of the algorithm. Thus, c is the analogue of the temperature in the physical annealing process. Note that the aforementioned probability decreases for increasing values of $C(j) - C(i)$ and for decreasing values of c, and cost-decreasing transitions are always accepted.

For a fixed value of c, the configurations that are consecutively visited by the algorithm can be seen as a Markov chain with transition matrix $P = P(c)$ given by (Aarts and Van Laarhoven 1985a, Lundy and Mees 1986, and Romeo and Sangiovanni-Vincentelli 1985)

$$P_{ij}(c) = \begin{cases} G_{ij}A_{ij}(c) & \text{if } j \neq i \\ 1 - \sum_{k=1}^{|\mathscr{R}|} G_{ik}A_{ik}(c) & \text{if } j = i, \end{cases} \tag{5}$$

where the generation probabilities G_{ij} are given by

$$G_{ij}(c) = \begin{cases} |\mathscr{N}(i)|^{-1} & \text{if } j \in \mathscr{N}(i) \\ 0 & \text{otherwise,} \end{cases} \tag{6}$$

and the acceptance probabilities A_{ij} by

$$A_{ij}(c) = \min\left\{1, \exp\left(\frac{-(C(j) - C(i))}{c}\right)\right\}. \tag{7}$$

The stationary distribution of this Markov chain exists and is given by (Aarts and Van Laarhoven 1985a, Lundy and Mees 1986, and Romeo and Sangiovanni-Vincentelli 1985):

$$q_i(c) = \frac{|\mathscr{N}(i)| A_{i_0 i}(c)}{\sum_{j \in \mathscr{R}} |\mathscr{N}(j)| A_{i_0 j}(c)} \tag{8}$$

for some $i_0 \in \mathscr{R}_{\text{opt}}$, where $\mathscr{R}_{\text{opt}}$ is the set of globally minimal configurations, provided the neighborhoods are such that for each pair of configurations (i, j) there is a finite sequence of transitions leading from i to j. The latter condition is equivalent to the requirement that the matrix G be irreducible. It can readily be shown that

$$\lim_{c \downarrow 0} q_i(c) = \begin{cases} |\mathscr{R}_{\text{opt}}|^{-1} & \text{if } i \in \mathscr{R}_{\text{opt}} \\ 0 & \text{otherwise.} \end{cases} \tag{9}$$

We recall that the stationary distribution of the Markov chain is defined as the probability distribution of the configurations after an infinite number of transitions. Thus, we conclude from (9) that the simulated annealing algorithm converges with probability 1 to a globally minimal configuration if the following conditions are satisfied:

- the sequence of values of the control parameter converges to 0;
- the Markov chains generated at each value of c are of infinite length; and
- the matrix G is irreducible.

Unfortunately, the neighborhood structure chosen for job shop scheduling in Section 3 is such that the corresponding matrix G is not irreducible. In that case, we can still prove asymptotic convergence provided the neighborhoods are such that for each configuration i there is a finite sequence of transitions leading from i to some configuration $i_0 \in \mathcal{R}_{opt}$ (Van Laarhoven 1988). To do so, we use the fact that in every chain the recurrent configurations can be divided uniquely into irreducible ergodic sets $\mathcal{S}_1$, $\mathcal{S}_2, \ldots, \mathcal{S}_T$. In addition to the ergodic sets there is a set $\mathcal{T}$ of transient configurations from which configurations in the ergodic sets can be reached (but not vice versa). Note that if the neighborhoods satisfy the aforementioned condition, then each $\mathcal{S}_t$ contains at least one globally minimal configuration.

Now consider the sequence of configurations constituting the Markov chain associated with $P(c)$. There are two possibilities: either the Markov chain starts in a transient configuration or it does not. In the latter case, the configurations constituting the Markov chain all belong to the same irreducible ergodic set $\mathcal{S}_t$ and we can prove asymptotic convergence as before, with $\mathcal{R}$ replaced by $\mathcal{S}_t$. On the other hand, if the Markov chain starts in a transient configuration, it will eventually "land" (Feller 1950) in an ergodic set $\mathcal{S}_t$, $t \in \{1, \ldots, T\}$, though it is not a priori known which one. The line of reasoning described above can then be applied again.

We can make the preceding arguments more precise by introducing the notion of a *stationary matrix Q*, whose elements q_{ij} are defined by

$$q_{ij} = \lim_{k \to \infty} \Pr\{X(k) = j \mid X(0) = i\}. \tag{10}$$

Using the results in Chapter 15, Sections 6–8 of Feller, we obtain

$$q_{ij} = \begin{cases} 0 & \text{if } j \in \mathcal{T} \text{ or } i \in \mathcal{S}_t, j \notin \mathcal{S}_t \\ & \quad \text{for some } t \in \{1, \ldots, T\}, \\[2mm] \dfrac{A_{i_0 j}(c)}{\sum_{l \in \mathcal{S}_t} A_{i_0 l}(c)} & \text{if } i, j \in \mathcal{S}_t \\ & \quad \text{for some } t \in \{1, \ldots, T\}, \\[2mm] x_{it} \dfrac{A_{i_0 j}(c)}{\sum_{l \in \mathcal{S}_t} A_{i_0 l}(c)} & \text{if } i \in \mathcal{T}, j \in \mathcal{S}_t \\ & \quad \text{for some } t \in \{1, \ldots, T\}, \end{cases} \tag{11}$$

where x_{it} is the probability that the Markov chain, starting from the transient configuration i, eventually reaches the ergodic set $\mathcal{S}_t$.

From (11) we obtain, for a recurrent configuration $j \in \mathcal{S}_t$,

$$0 \leq \lim_{k \to \infty} \Pr\{X(k) = j\} = \sum_{i \in \mathcal{R}} \Pr\{X(0) = i\} \cdot q_{ij}$$

$$= \left(\sum_{i \in \mathcal{T}} \Pr\{X(0) = i\} \cdot x_{it} + \sum_{i \in \mathcal{S}_t} \Pr\{X(0) = i\} \right)$$

$$\cdot \frac{A_{i_0 j}(c)}{\sum_{l \in \mathcal{S}_t} A_{i_0 l}(c)}$$

$$\leq \frac{A_{i_0 j}(c)}{\sum_{l \in \mathcal{S}_t} A_{i_0 l}(c)}. \tag{12}$$

Using (7) we find

$$\lim_{c \downarrow 0} \frac{A_{i_0 j}(c)}{\sum_{l \in \mathcal{S}_t} A_{i_0 l}(c)} = 0 \tag{13}$$

if $j \in \mathcal{S}_t, j \notin \mathcal{R}_{opt}$. Consequently,

$$\lim_{c \downarrow 0}(\lim_{k \to \infty} \Pr\{X(k) = j\}) = 0$$

for any transient or nonglobally minimal recurrent configuration j. In other words

$$\lim_{c \downarrow 0}\left(\lim_{k \to \infty} \Pr\{X(k) \in \mathcal{R}'_{opt}\} \right) = 1, \tag{14}$$

where $\mathcal{R}'_{opt}$ denotes the nonempty set of globally minimal recurrent configurations.

Some of the conditions for asymptotic convergence, as, for instance, the infinite length of Markov chains, cannot be met in practice. In any finite-time implementation, we therefore have to make a choice with respect to each of the following parameters:

- the length of the Markov chains,
- the initial value of the control parameter,
- the decrement rule of the control parameter,
- the final value of the control parameter.

Such a choice is usually referred to as a *cooling schedule* or an *annealing schedule*. The results in this paper have been obtained by an implementation of simulated annealing that employs the cooling schedule derived by Aarts and Van Laarhoven (1985a, b). This is a three-parameter schedule: The parameters χ_0 and ϵ_s determine the initial and final values of the control parameter, respectively, whereas the decrement rule depends on a parameter δ in the following way:

$$c_{k+1} = \frac{c_k}{1 + [c_k \cdot \ln(1 + \delta)/3\sigma_k]} \tag{15}$$

where c_k is the value of the control parameter for the kth Markov chain and σ_k is the standard deviation of the cost values of the configurations obtained by generating the kth Markov chain.

The decrement of c_k is such that the stationary distributions of two succeeding Markov chains approximately satisfy (Aarts and Van Laarhoven 1985a)

$$\frac{1}{1 + \delta} < \frac{q_i(c_k)}{q_i(c_{k+1})} < 1 + \delta,$$

$$k = 1, 2, \ldots \quad \text{for all } i \in \mathscr{R}. \quad (16)$$

Thus, for small values of δ, the stationary distributions of succeeding Markov chains are "close" to each other and we therefore argue that, after decreasing c_k to c_{k+1}, a small number of transitions suffices to let the probability distribution of the configurations approach the new stationary distribution $q(c_{k+1})$. Note that small values of δ correspond to a slow decrement of the control parameter.

Finally, we choose the length of each Markov chain, L_k, equal to the size of the largest neighborhood, i.e.,

$$L_k = \max_{i \in \mathscr{R}} |\mathscr{N}(i)|, \quad k = 1, 2, \ldots. \quad (17)$$

We have applied simulated annealing based on this cooling schedule to many problems in combinatorial optimization (see, for example, Van Laarhoven and Aarts) and have found it extremely robust in that the final results are typically within 2% of the global minimum, when δ is chosen sufficiently low (0.1 or smaller).

Under some mild assumptions, it is possible to show that the aforementioned cooling schedule leads to a time-complexity of the simulated annealing algorithm given by $\mathscr{O}(\tau L \ln |\mathscr{R}|)$, where τ is the time involved in the generation and (possible) acceptance of a transition and L is the size of the largest neighborhood (the length of the Markov chains) (Aaarts and Van Laarhoven 1985a). If one works out this bound for a particular combinatorial optimization problem, it is usually polynomial in the size of the problem. In those cases, we have a polynomial-time approximation algorithm. Such a result with respect to the *efficiency* of the algorithm is only worthwhile in combination with results on its *effectiveness*, viz. on the difference in cost between solutions returned by the algorithm and globally minimal ones. From a theoretical point of view, very little is known about the effectiveness of simulated annealing, but there are many empirical results; see for instance the extensive computational experiments of Johnson et al. (1989). For the job shop scheduling problem, we present an empirical analysis of the effectiveness and efficiency of simulated annealing in Section 4, but first the application of simulated annealing to the job shop scheduling problem is discussed in more detail.

3. SIMULATED ANNEALING AND JOB SHOP SCHEDULING

We recall from the previous section that in order to apply simulated annealing to any combinatorial optimization problem, we need a precise definition of configurations, a cost function and a neighborhood structure. Furthermore, to prove asymptotic convergence we must show that the neighborhood structure is such that for an arbitrary configuration i there exists at least one globally minimal configuration $i_0 \in \mathscr{R}_{opt}$ that can be reached from i in a finite number of transitions. Hereinafter, we discuss these items in more detail.

3.1. Configurations

We recall from Section 1 that we can solve the job shop scheduling problem by considering sets of machine permutations and by determining, for such a set of permutations, the longest path in the digraph which results from giving the edges in the disjunctive graph the orientations determined by the permutations. We therefore define a configuration i of the problem as a set $\Pi_i = \{\pi_{i1}, \ldots, \pi_{im}\}$ of machine permutations. π_{ik} is to be interpreted as the order in which the operations on machine k are processed: If $M_v = k$ for some $v \in \mathscr{O}$, then $\pi_{ik}(v)$ denotes the operation following v on machine k. Consequently, the number of configurations is given by $\Pi_{k=1}^m m_k!$, where m_k is the number of operations to be processed by machine k ($m_k = |\{v \in \mathscr{O} | M_v = k\}|$).

3.2. Cost Function

For each configuration i we define the following two digraphs:

$D_i = (V, A \cup E_i)$, where

$E_i = \{(v, w) | \{v, w\} \in E \text{ and } \pi_{ik}(v) = w$

$$\text{for some } k \in \mathscr{M}\}. \quad (18)$$

$\bar{D}_i = (V, A \cup \bar{E}_i)$, where

$\bar{E}_i = \{(v, w) | \{v, w\} \in E \text{ and } \pi_{ik}^l(v) = w$

$$\text{for some } k \in \mathscr{M}, 1 \leq l \leq m_k - 1\}. \quad (19)$$

In other words, $\bar{D}_i$ is the digraph obtained from the disjunctive graph by giving the edges in E the orientations resulting from Π_i; the digraph D_i can be obtained from $\bar{D}_i$ by taking only those arcs from $\bar{E}_i$ that connect successive operations on the same

machine. It is well known that the longest paths in D_i and $\bar{D}_i$ have equal length. Thus, the cost of a configuration i can be found by determining the length of a longest path from 0 to $N + 1$ in D_i. To compute such a cost, we use a simple *labeling algorithm*, based on Bellman's equations (Bellman 1958), for solving the longest-path problem in a digraph. The time-complexity of this algorithm is proportional to the number of arcs in the graph. In our case, this number equals $|A| + |E_i| = (N + n) + (N - m)$; accordingly, the labeling algorithm takes $\mathcal{O}(N)$ time to compute the cost of a configuration.

3.3. Neighborhood Structure

A transition is generated by choosing vertices v and w, such that:

1. v and w are successive operations on some machine k;
2. $(v, w) \in E_i$ is a critical arc, i.e., (v, w) is on a longest path in D_i;

and reversing the order in which v and w are processed on machine k. Thus, in the digraph D_i such a transition results in reversing the arc connecting v and w and replacing the arcs (u, v) and (w, x) by (u, w) and (v, x), respectively, where $u = \pi_{ik}^{-1}(v)$ and $x = \pi_{ik}(w)$. Our choice is motivated by two facts:

- Reversing a critical arc in a diagraph D_i can never lead to a cyclic diagraph D_j (see Lemma 2).
- If the reversal of a noncritical arc in D_i leads to an acyclic graph D_j, a longest path q in D_j cannot be shorter than a longest path p in D_i (because D_j still contains the path p).

Thus, we exclude beforehand some noncost-decreasing transitions and, in addition, all transitions that might result in a cyclic digraph. Consequently, the neighborhood structure is such that the algorithm visits only digraphs corresponding to feasible solutions.

The neighborhood of a configuration i is thus given by the set of acyclic diagraphs that can be obtained by reversing a critical arc belonging to E_i in the graph D_i. Consequently, $|\mathcal{N}(i)| < \sum_{k=1}^{m} (m_k - 1) = N - m$.

3.4. Asymptotic Convergence

It is not difficult to construct a problem instance containing pairs of configurations (i, j) for which there is no finite sequence of transitions leading from i to j (Van Laarhoven). Thus, to prove asymptotic convergence, we must show that for each configuration i there is a finite sequence of transitions leading from i

to some globally minimal configuration. In order to do so, we need two lemmas.

Lemma 1. *Consider an arbitrary configuration i and an arbitrary global minimum $i_0 \in \mathcal{R}_{\text{opt}}$. If $i \notin \mathcal{R}_{\text{opt}}$, then the set $K_i(i_0)$ defined by*

$$K_i(i_0) = \{e = (v, w) \in E_i \mid e$$
$$\text{is critical} \wedge (w, v) \in \bar{E}_{i_0}\} \quad (20)$$

is not empty.

Proof. The proof consists of two parts: First, we show that E_i always contains critical arcs, unless $i \in \mathcal{R}_{\text{opt}}$; next that there are always critical arcs in E_i that do not belong to $\bar{E}_{i_0}$ unless again $i \in \mathcal{R}_{\text{opt}}$.

1. Suppose that E_i contains no critical arcs, then all critical arcs belong to A. Consequently, a longest path consists of arcs connecting vertices corresponding to operations of the same job; accordingly, its length is given by the total processing time of that job. But this is a lower bound to the length of a longest path in any digraph D_j, hence $i \in \mathcal{R}_{\text{opt}}$.
2. Suppose that for all critical arcs e in E_i, we have $e \in \bar{E}_{i_0}$. We then know that any longest path p in D_i is also a path q in $\bar{D}_{i_0}$. The length of a longest path r in $\bar{D}_{i_0}$ is also the length of a longest path in D_{i_0} and because $i_0 \in \mathcal{R}_{\text{opt}}$, we have $length(r) \leqslant length(p)$. But by definition $length(r) \geqslant length(q) = length(p)$. Consequently, $length(p) = length(r)$ and $i \in \mathcal{R}_{\text{opt}}$.

Lemma 2. *Suppose that $e = (v, w) \in E_i$ is a critical arc of an acyclic digraph D_i. Let D_j be the digraph obtained from D_i by reversing the arc e in E_i. Then D_j is also acyclic.*

Proof. Suppose that D_j is cyclic. Because D_i is acyclic, the arc (w, v) is part of the cycle in D_j. Consequently, there is a path $(v, x, y, \ldots, w)$ in D_j. But this path can also be found in D_i and is clearly a longer path from v to w than the arc (v, w). This contradicts the assumption that (v, w) is on a longest path in D_i. Hence, D_j is acyclic.

Given a configuration $i_0 \in \mathcal{R}_{\text{opt}}$, we define two sets for an arbitrary configuration i:

$$M_i(i_0) = \{e = (v, w) \in E_i \mid (w, v) \in \bar{E}_{i_0}\} \quad (21)$$

$$\bar{M}_i(i_0) = \{e = (v, w) \in \bar{E}_i \mid (w, v) \in \bar{E}_{i_0}\}. \quad (22)$$

In view of Section 2, the following theorem now ensures asymptotic convergence in probability to a globally minimal configuration.

Theorem 1. *For each configuration $i \notin \mathcal{R}_{opt}$ it is possible to construct a finite sequence of transitions leading from i to a globally minimal configuration.*

Proof. We choose an arbitrary configuration $i_0 \in \mathcal{R}_{opt}$ and construct a sequence of configurations $\{\lambda_0, \lambda_1, \ldots\}$ as follows:

1. $\lambda_0 = i$.
2. λ_{k+1} is obtained from λ_k by reversing an arc $e \in K_{\lambda_k}(i_0)$ in E_{λ_k}. According to Lemma 2, this can be done without creating a cycle in $D_{\lambda_{k+1}}$. Furthermore, this operation is of the aforementioned type of transition.

It is easy to see that if $|\bar{M}_{\lambda_k}(i_0)| > 0$ then

$$|\bar{M}_{\lambda_{k+1}}(i_0)| = |\bar{M}_{\lambda_k}(i_0)| - 1. \tag{23}$$

Hence, for $k = |\bar{M}_i(i_0)|$, $|\bar{M}_{\lambda_k}(i_0)| = 0$. Using $K_i(i_0) \subseteq M_i(i_0) \subseteq \bar{M}_i(i_0)$, we find $K_{\lambda_k}(i_0) = \varnothing$ for $k = |\bar{M}_i(i_0)|$. According to Lemma 1, this implies $\lambda_k \in \mathcal{R}_{opt}$.

4. COMPUTATIONAL RESULTS

We have analyzed the finite-time behavior of the simulated annealing algorithm empirically by running the algorithm on a number of instances of the job shop scheduling problem, varying in size from six jobs on six machines to thirty jobs on ten machines. For all instances, the number of operations of each job equals the number of machines and each job has precisely one operation on each machine. In that case, the number of configurations of each instance is given by $(n!)^m$, the labeling algorithm takes $\mathcal{O}(nm)$ time to compute the cost of a configuration, and the size of the neighborhood of a configuration is bounded by $m(n - 1)$.

FIS1, FIS2 and FIS3 (Table I) are three problem instances due to Fisher and Thompson (1963), the forty instances in Table II are due to Lawrence (1984). FIS2 is a notorious 10-job, 10-machine instance that has defied solution to optimality for more than twenty years. A couple of years ago, a solution with cost 930 was found after 1 hour 47 minutes of running time, and no improvement was found after 9 hours 6 minutes (Lageweg 1984). This cost value was only recently proved to be globally minimal by Carlier and Pinson (1989). For FIS1, FIS2 and FIS3, the processing times of the operations are randomly drawn and range from 1 to 10 (FIS1) or to 99 (FIS2 and FIS3) units of time. The sequence of machines for each job is such that lower-numbered machines tend to be used for earlier operations. For the Lawrence instances, processing times are drawn from a uniform distribution on the interval [5, 99]; the sequence of machines for each job is random.

The performance of simulated annealing on these instances is reported in Table I for the Fisher–Thompson instances, and in Table II for the Lawrence instances. The averages in these tables are computed from five solutions, obtained by running the algorithm, controlled by the cooling schedule described in Section 2, five times on each instance and recording the best configuration encountered during each run (this need not necessarily be the final configuration). The probabilistic nature of the algorithm makes it necessary to carry out multiple runs on the same problem instance in order to get meaningful results.

All results are obtained with the parameters χ_0 and ϵ_s set to 0.95 and 10^{-6}, respectively, and for different values of the distance parameter δ. Running times are CPU times on a VAX-785.

From Tables I and II we can make the following observations:

1. The quality of the average best solution returned by the algorithm improves considerably when δ is decreased. This is in accordance with the theory underlying the employed cooling schedule: smaller values of δ correspond to a better approximation of the asymptotic behavior (Aarts and Van Laarhoven 1985a). Furthermore, the difference between the average best solution and a globally minimal one does not deteriorate significantly with increasing problem size. For the FIS2 instance, the five best solutions obtained with $\delta = 10^{-4}$ have cost values of 930 (twice), 934, 935, and 938, respectively. Thus, a globally minimal solution is found 2 out of 5 times, which is quite a remarkable result, considering the notoriety of this instance.
2. As for running times, the bound for the running time given in Section 2 is $\mathcal{O}((nm)^3 \ln n)$ ($L = \mathcal{O}(nm)$, $|\mathcal{R}| = \mathcal{O}((n!)^m)$ and $\tau = \mathcal{O}(nm)$). Thus, for fixed m the bound is $\mathcal{O}(n^3 \ln n)$, and for fixed n it is $\mathcal{O}(m^3)$. For the A, B and C instances in Table II, for which m is constant, the average running time $\bar{t}$ for $\delta = 0.01$ is approximately given by $\bar{t} = t_0 \cdot n^{2.215} \cdot \ln n$, for some constant t_0 ($\chi^2 = 1.00$); for the G, B and I instances, for which n is constant, the average running time for $\delta = 0.01$ is approximately given by $\bar{t} = t_1 \cdot m^{2.406}$, for some constant t_1 ($\chi^2 = 1.00$). Thus, the observed running times are in good accordance with the bound given in Section 2.

Table I
Results for Problem Instances of Fisher and Thompson (1963)[a]

| Problem | δ | Simulated Annealing | | | | | | | ABZ | | MSS | |
		$\bar{C}$	σ_C	%	$\bar{t}$	σ_t	C_{best}	C_A	t_A	C_M	t_M
6 machines, 6 jobs											
FIS1	10^{-1}	56.0	1.3	1.82	8	1	55*	55*	1	—	—
	10^{-2}	55.0*	0.0	0.00	52	8	55*				
10 machines, 10 jobs											
FIS2	10^{-1}	1,039.6	15.1	11.78	113	13	1,028	930*	426	946	494
	10^{-2}	985.8	22.1	6.00	779	61	951				
	10^{-3}	942.4	4.5	1.33	5,945	180	937				
	10^{-4}	933.4	3.1	0.37	57,772	2,364	930*				
20 machines, 5 jobs											
FIS3	10^{-1}	1,354.2	26.5	16.24	123	13	1,325	1,178	40	—	—
	10^{-2}	1,229.0	33.6	5.49	848	93	1,184				
	10^{-3}	1,187.0	18.7	1.89	6,840	389	1,173				
	10^{-4}	1,173.8	5.2	0.76	62,759	7,805	1,165*				

| | Iterative Improvement | | | | | | |
	#	$\bar{C}$	σ_C	%	$\bar{t}$	σ_t	C_{best}
FIS1	803.2	55.4	0.8	0.73	52	0	55*
FIS2	9,441.2	1,018.2	9.1	9.48	5,945	0	1,006
FIS3	5,221.0	1,331.4	9.5	14.28	6,841	0	1,319

[a] Simulated annealing:

δ : the distance parameter (smaller δ-values imply slower cooling);

$\bar{C}, \sigma_C$: the average cost and standard deviation of best solution over five runs;

$\bar{t}, \sigma_t$: the average running time and standard deviation over five runs (in seconds);

C_{best} : the best cost found over five runs;

% : the percentage of $\bar{C}$ over optimal cost.

Iterative improvement:

\# : the average number of local minima over five macroruns (see text);

$\bar{C}, \sigma_C$: the average cost and standard deviation of best solution over five macroruns;

$\bar{t}, \sigma_t$: the average running time and standard deviation over five macroruns (in seconds);

C_{best} : the best cost found over five macroruns;

% : the percentage of $\bar{C}$ over optimal cost.

Adams, Balas and Zawack (1988):

C_A : the best cost found;

t_A : running time (in seconds).

Matsuo, Suh and Sullivan (1988):

C_M : the best cost found;

t_M : running time (in seconds);

* : provably optimal cost.

In the remainder of this section we compare the results obtained with our method with results obtained with three other methods, viz.

- time-equivalent iterative improvement,
- the shifting bottleneck procedure (Adams, Balas and Zawack), and
- controlled search simulated annealing (Matsuo, Suh and Sullivan).

4.1. Time-Equivalent Iterative Improvement

Table I also contains results obtained by repeated execution of a time-equivalent iterative improvement algorithm based on the same neighborhood structure as our simulated annealing algorithm. It is based on repeated execution of iterative improvement. The averages for the time-equivalent iterative improve-

ment are obtained from five macroruns. Each macrorun consists of repeated execution of the iterative improvement algorithm for a large number of randomly generated initial configurations and thus yields a large number of local minima. Execution of each macrorun is terminated as soon as the running time is equal to the running time of an average run of simulated annealing applied to the same problem instance with the distance parameter δ set to 10^{-3} (10^{-2} for **FIS1**); $\bar{C}$ is the average of the best cost value found during each macrorun.

We observe that repeated execution of iterative improvement is easily outperformed by simulated annealing for the two larger problems. The difference is significant: for **FIS3**, for instance, the average best solution obtained by simulated annealing is almost 11% better in cost than the one obtained by repeated execution of iterative improvement.

Table II
Results for Problem Instances of Lawrence (1984)[a]

Problem	δ	$\bar{C}$	σ_C	$\bar{t}$	σ_t	C_{best}	C_A	t_A	C_M	t_M
			Simulated Annealing				ABZ		MSS	
10 machines, 10 jobs										
A1	1.0	1023.4	30.6	26	1.5	991	978	120	959	78
	0.1	981.0	17.3	110	17.1	956				
	0.01	966.2	10.1	686	83.3	956				
A2	1.0	861.0	41.2	23	3.7	797	787	96	784	47
	0.1	792.4	6.2	112	7.0	784				
	0.01	787.8	1.6	720	109.0	785				
A3	1.0	902.6	30.9	23	1.6	870	859	112	848	53
	0.1	872.2	12.4	112	22.1	861				
	0.01	861.2	0.4	673	69.0	861				
A4	1.0	950.0	54.5	24	5.3	904	860	120	842	58
	0.1	881.4	6.9	97	20.4	874				
	0.01	853.4	4.6	830	85.4	848				
A5	1.0	1021.6	26.2	30	1.9	994	914	144	907	50
	0.1	927.6	18.9	86	7.9	907				
	0.01	908.4	4.2	667	126.9	902				
10 machines, 15 jobs										
B1	1.0	1176.2	37.8	69	6.7	1133	1084	181	1071	103
	0.1	1115.2	23.9	299	50.9	1085				
	0.01	1067.6	3.7	1991	341.1	1063				
B2	1.0	1125.6	35.6	65	3.6	1094	944	210	927	92
	0.1	977.4	19.5	307	36.5	963				
	0.01	944.2	4.7	2163	154.6	938				
B3	1.0	1155.8	64.2	63	5.6	1056	1032*	113	1032*	10
	0.1	1051.0	24.6	275	35.8	1032*				
	0.01	1032.0*	0.0	2093	89.7	1032*				
B4	1.0	1101.0	53.5	71	5.0	1032	976	217	973	100
	0.1	977.6	8.1	252	28.5	968				
	0.01	966.6	8.7	2098	406.0	952				
B5	1.0	1114.6	9.1	77	16.9	1103	1017	215	991	90
	0.1	1035.4	10.6	283	44.3	1017				
	0.01	1004.4	14.4	2133	374.5	992				
10 machines, 20 jobs										
C1	1.0	1397.0	69.1	139	16.0	1311	1224	372	1218*	27
	0.1	1268.0	9.7	555	81.7	1252				
	0.01	1219.0	2.0	4342	597.8	1218*				
C2	1.0	1434.2	40.0	139	6.4	1390	1291	419	1274	143
	0.1	1311.6	12.7	651	82.9	1295				
	0.01	1273.6	5.2	4535	392.0	1269				
C3	1.0	1414.6	57.8	135	7.4	1335	1250	451	1216	153
	0.1	1280.2	23.6	614	83.3	1246				
	0.01	1244.8	15.4	4354	349.8	1224				
C4	1.0	1387.4	47.0	138	14.1	1307	1239	446	1196	134
	0.1	1260.4	35.4	581	24.0	1203				
	0.01	1226.4	6.5	4408	450.9	1218				
C5	1.0	1539.2	44.2	145	20.6	1492	1355*	276	1355*	4
	0.1	1393.6	9.6	605	84.4	1381				
	0.01	1355.0*	0.0	3956	428.2	1355*				

Table II—*Continued*
Results for Problem Instances of Lawrence (1984)

Problem	δ	Simulated Annealing					ABZ		MSS	
		$\bar{C}$	σ_C	$\bar{t}$	σ_t	C_{best}	C_A	t_A	C_M	t_M
10 machines, 30 jobs										
D1	1.0	1882.2	39.3	442	79.3	1821	1784*	19	—	—
	0.1	1784.0*	0.0	1517	58.1	1784*				
D2	1.0	1921.4	35.3	492	66.2	1868	1850*	15	—	—
	0.1	1850.0*	0.0	1752	124.6	1850*				
D3	1.0	1761.8	12.2	433	40.4	1740	1719*	14	—	—
	0.1	1726.6	15.2	1880	130.8	1719*				
D4	1.0	1816.4	27.7	470	31.2	1788	1721*	11	—	—
	0.1	1775.6	38.4	1886	232.4	1721*				
D5	1.0	2011.2	81.3	434	34.6	1888*	1888*	11	—	—
	0.1	1890.0	4.0	1668	107.9	1888*				
5 machines, 10 jobs										
F1	1.0	707.0	32.2	6	1.0	666*	666*	1	—	—
	0.1	666.0*	0.0	20	3.5	666*				
	0.01	666.0*	0.0	123	15.3	666*				
F2	1.0	719.0	20.0	6	1.0	685	669	6	655*	2
	0.1	671.0	11.1	24	2.5	655*				
	0.01	663.0	4.9	117	19.0	655*				
F3	1.0	689.6	22.4	5	0.9	664	605	32	597	17
	0.1	635.6	9.5	24	3.8	626				
	0.01	617.6	8.5	129	12.6	606				
F4	1.0	665.4	56.9	6	0.9	608	593	23	590	17
	0.1	617.2	20.5	21	5.2	594				
	0.01	593.8	2.1	121	15.9	590				
F5	1.0	594.4	2.8	5	0.6	593*	593*	0	—	—
	0.1	593.0*	0.0	19	4.2	593*				
	0.01	593.0*	0.0	118	15.3	593*				
5 machines, 15 jobs										
G1	1.0	937.2	13.7	16	2.8	926*	926*	1	—	—
	0.1	926.0*	0.0	52	5.8	926*				
	0.01	926.0*	0.0	286	32.1	926*				
G2	1.0	948.6	44.1	15	1.6	911	890*	1	—	—
	0.1	900.6	8.5	66	15.2	890*				
	0.01	890.0*	0.0	376	48.3	890*				
G3	1.0	905.8	34.2	16	0.4	863*	863*	2	863*	1
	0.1	863.0*	0.0	55	7.3	863*				
	0.01	863.0*	0.0	292	40.8	863*				
G4	1.0	965.2	20.0	13	1.0	951*	951*	0	—	—
	0.1	951.0*	0.0	47	5.9	951*				
	0.01	951.0*	0.0	283	25.9	951*				
G5	1.0	958.0*	0.0	14	1.6	958*	958*	0	—	—
	0.1	958.0*	0.0	45	2.0	958*				
	0.01	958.0*	0.0	243	42.3	958*				

Table II—*Continued*
Results for Problem Instances of Lawrence (1984)[a]

Problem	δ	$\bar{C}$	σ_C	$\bar{t}$	σ_t	C_{best}	C_A	t_A	C_M	t_M
			Simulated Annealing				ABZ		MSS	
5 machines, 20 jobs										
H1	1.0	1229.6	14.7	32	3.9	1222*	1222*	1	—	—
	0.1	1222.0*	0.0	108	17.2	1222*				
	0.01	1222.0*	0.0	627	18.4	1222*				
H2	1.0	1042.8	7.6	34	3.9	1039*	1039*	0	—	—
	0.1	1061.2	44.4	116	11.9	1039*				
	0.01	1039.0*	0.0	655	30.7	1039*				
H3	1.0	1154.6	9.2	32	2.5	1150*	1150*	1	—	—
	0.1	1150.0*	0.0	118	18.0	1150*				
	0.01	1150.0*	0.0	564	85.9	1150*				
H4	1.0	1292.0*	0.0	27	1.7	1292*	1292*	0	—	—
	0.1	1292.0*	0.0	93	20.6	1292*				
	0.01	1292.0*	0.0	462	21.8	1292*				
H5	1.0	1299.8	77.4	34	5.3	1207*	1207*	2	—	—
	0.1	1252.5	18.8	126	16.2	1233				
	0.01	1207.0*	0.0	736	26.3	1207*				
15 machines, 15 jobs										
I1	1.0	1487.6	40.4	152	6.4	1450	1305	268	1292	312
	0.1	1343.2	30.2	785	80.6	1297				
	0.01	1300.0	7.8	5346	399.8	1293				
I2	1.0	1580.2	38.3	173	12.1	1523	1423	419	1435	289
	0.1	1479.4	28.8	757	94.6	1457				
	0.01	1442.4	5.7	5287	688.5	1433				
I3	1.0	1422.2	28.3	173	25.3	1376	1255	540	1231	336
	0.1	1303.4	30.5	713	90.8	1263				
	0.01	1227.2	8.2	5480	614.8	1215				
I4	1.0	1408.6	44.4	186	24.6	1348	1273	335	1251	330
	0.1	1305.4	27.5	673	75.1	1264				
	0.01	1258.2	5.2	5766	800.3	1248				
I5	1.0	1399.8	60.2	162	8.8	1318	1269	450	1235	308
	0.1	1282.2	15.5	745	68.4	1254				
	0.01	1247.4	9.9	5373	1066.4	1234				

[a] The legend for this table appears on the bottom of Table I.

4.2. The Shifting Bottleneck Procedure

Tables I and II also contain for each instance the cost value of the best solution obtained by Adams, Balas and Zawack with their shifting bottleneck procedure. Most values are obtained by a second heuristic, which embeds the aforementioned sliding bottleneck procedure and proceeds by partial enumeration of the solution space. The values for the instances F1, F5, G3 as well as for the D and H instances are obtained by the sliding bottleneck procedure only. The corresponding running times are obtained by halving the running times in Adams, Balas and Zawack, since these correspond to a VAX-780. Adams, Balas and Zawack show their approach to be superior to a number of approaches based on priority dispatching rules: The typical improvement is reported to be between 4% and 10%.

Comparison of simulated annealing and the shifting bottleneck procedure leads to the following observations:

1. For those instances for which Adams, Balas and Zawack do not find a globally minimal solution (mainly the A, B, C and I instances in Table II), the running times of simulated annealing with $\delta = 0.1$ and of the heuristic of Adams, Balas and Zawack are of the same order of magnitude. In this case, the best solution found by Adams, Balas and Zawack is considerably better than the average

best solution returned by simulated annealing and as good as the best solution found in five runs of simulated annealing.

Putting $\delta = 0.01$ makes simulated annealing much slower than the heuristic of Adams, Balas and Zawack, but now their best solution is slightly worse than the average best solution of simulated annealing and considerably worse than the best solution in five runs of simulated annealing (the typical improvement is between 1 and 3%).

2. For the instances for which the heuristic of Adams, Balas and Zawack finds a globally minimal solution, it outperforms simulated annealing: The latter algorithm also finds global minima, but takes much more time to do so.

Admittedly, the performance of the shifting bottleneck heuristic is liable to improve if it is allowed more time. Nevertheless, the results in Tables I and II indicate that simulated annealing is a promising approach to job shop scheduling, as well as a robust one (cf. the small difference between C and C_{best} in Tables I and II for $\delta = 0.01$) and certainly superior to traditional approaches, such as procedures based on priority dispatching rules.

4.3. Controlled Search Simulated Annealing

Finally, Tables I and II contain for a number of instances the cost value of the best solution obtained by Matsuo, Suh and Sullivan with their controlled search simulated annealing method. They use neighborhoods consisting of schedules that can be obtained by several types of (multi)adjacency interchanges of operations that are critical for determining the makespan. Briefly, these neighborhoods are obtained by augmenting the relatively simple neighborhoods used in our approach by adding better schedules that are found by exploring the structure of the problem at hand. Evidently, this makes the approach less general than ours. As the tables show, this augmentation enhances the efficiency of the algorithm but not its effectiveness; the quality of the final solution remains roughly the same.

The running times for this approach, given in Tables I and II, are again obtained by halving the times given by Matsuo, Suh and Sullivan, since they also used a VAX-780 computer. Comparison of our simulated annealing method with controlled search simulated annealing yields the following conclusions.

1. For those instances for which Matsuo, Suh and Sullivan do not find an optimum, our approach finds, on the average, solutions of the same quality but at the cost of larger computational efforts.

2. For those instances for which an optimum was found, the controlled search simulated annealing method finds it in remarkably smaller running times, even when compared to the running times used by the shifting bottleneck procedure.

5. CONCLUSION

We have discussed a new approach to job shop scheduling based on a randomized version of iterative improvement. The probabilistic element of the algorithm (the acceptance of cost-increasing transitions with a nonzero probability) makes simulated annealing a significantly better approach than the classical iterative improvement approach on which it is based. The difference is especially pronounced for large problem instances.

For a number of well-studied problems in combinatorial optimization a comparison of simulated annealing and tailored heuristics usually leads to the conclusion that tailored algorithms are more efficient and more effective than simulated annealing: They find better solutions in less time (see, for example, Van Laarhoven and Aarts, and Johnson et al.). Interestingly, this does not seem to be entirely the case for job shop scheduling: Simulated annealing has a potential of finding shorter makespans than the tailored heuristic of Adams, Balas and Zawack, at the cost of large running times. In other words, tailoring the algorithm toward the combinatorial structure of the problem does not yield a more effective, merely a more efficient algorithm. This observation is confirmed by the results of Matsuo, Suh and Sullivan, who describe a simulated annealing-based approach to job shop scheduling, employing a more problem specific neighborhood structure. Their approach also leads to a more efficient algorithm, but it again produces results of the same quality.

We consider the disadvantage of large running times to be compensated for by the simplicity of the algorithm, by its ease of implementation, by the fact that it requires no deep insight into the combinatorial structure of the problem, and, of course, by the high quality of the solutions it returns.

REFERENCES

AARTS, E. H. L., AND P. J. M. VAN LAARHOVEN. 1985a. Statistical Cooling: A General Approach to Combinatorial Optimization Problems. *Philips J. Res.* **40**, 193–226.

AARTS, E. H. L., AND P. J. M. VAN LAARHOVEN. 1985b. A New Polynomial Time Cooling Schedule. In

Proceedings IEEE International Conference on Computer-Aided Design, Santa Clara, Calif., 206–208 (November).

ADAMS, J., E. BALAS AND D. ZAWACK. 1988. The Shifting Bottleneck Procedure for Job Shop Scheduling. *Mgmt. Sci.* **34**, 391–401.

BELLMAN, R. E. 1958. On a Routing Problem. *Quart. Appl. Math.* **16**, 87–90.

CARLIER, J., AND E. PINSON. 1989. An Algorithm for Solving the Job-Shop Problem. *Mgmt. Sci.* **35**, 164–176.

ČERNY, V. 1985. Thermodynamical Approach to the Traveling Salesman Problem: An Efficient Simulation Algorithm. *J. Opt. Theory Appl.* **45**, 41–51.

COFFMAN, E. G., JR. (ed.). 1976. *Computer and Job-Shop Scheduling Theory.* John Wiley, New York.

FELLER, W. 1950. *An Introduction to Probability Theory and Applications*, Vol. 1. John Wiley, New York.

FISHER, H., AND G. L. THOMPSON. 1963. Probabilistic Learning Combinations of Local Job-Shop Scheduling Rules. In *Industrial Scheduling*. J. F. Muth and G. L. Thompson (eds.). Prentice Hall, Englewood Cliffs, N.J., 225–251.

FRENCH, S. 1982. *Sequencing and Scheduling: An Introduction to the Mathematics of the Job-Shop.* Horwood, Chichester, U.K.

JOHNSON, D. S., C. R. ARAGON, L. A. McGEOCH AND C. SCHEVON. 1989. Optimization by Simulated Annealing: An Experimental Evaluation (Part I). *Opns. Res.* **37**, 865–892.

KIRKPATRICK, S., C. D. GELATT, JR. AND M. P. VECCHI. 1983. Optimization by Simulated Annealing. *Science* **220**, 671–680.

LAGEWEG, B. J. 1984. Private Communication.

LAGEWEG, B. J., J. K. LENSTRA AND A. H. G. RINNOOY KAN. 1977. Job-Shop Scheduling by Implicit Enumeration. *Mgmt. Sci.* **24**, 441–450.

LAWLER, E. L., J. K. LENSTRA AND A. H. G. RINNOOY KAN. 1982. Recent Developments in Deterministic Sequencing and Scheduling: A Survey. In *Deterministic and Stochastic Scheduling*. M. A. H. Dempster, J. K. Lenstra and A. H. G. Rinnooy Kan (eds.). Reidel, Dordrecht, The Netherlands, 35–73.

LAWRENCE, S. 1984. Resource Constrained Project Scheduling: An Experimental Investigation of Heuristic Scheduling Techniques (Supplement). Graduate School of Industrial Administration, Carnegie Mellon University, Pittsburgh.

LUNDY, M., AND A. MEES. 1986. Convergence of an Annealing Algorithm. *Math. Prog.* **34**, 111–124.

MATSUO, H., C. J. SUH AND R. S. SULLIVAN. 1988. A Controlled Search Simulated Annealing Method for the General Jobshop Scheduling Problem. Working Paper 03-04-88, Department of Management, The University of Texas at Austin.

PAPADIMITRIOU, C. H., AND K. STEIGLITZ. 1982. *Combinatorial Optimization: Algorithms and Complexity.* Prentice-Hall, Englewood Cliffs, N.J.

ROMEO, F., AND A. L. SANGIOVANNI-VINCENTELLI. 1985. Probabilistic Hill Climbing Algorithms: Properties and Applications. In *Proceedings 1985 Chapel Hill Conference on VLSI*, Chapel Hill, N.C., 393–417 (May).

ROY, B., AND B. SAUSSMANN. 1964. Les Problèmes d'Ordonnancement avec Constraints Disjonctives. *Note DS No. 9 bis*, SEMA, Paris.

VAN LAARHOVEN, P. J. M. 1988. Theoretical and Computational Aspects of Simulated Annealing. Ph.D. Thesis, Erasmus University, Rotterdam.

VAN LAARHOVEN, P. J. M., AND E. H. L. AARTS. 1987. *Simulated Annealing: Theory and Applications.* Reidel, Dordrecht, The Netherlands.